Untitled8

May 31, 2024

[71]:
```
#importing the needed libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

[2]:
```
#LOADING THE DATASET
cancelled_3= pd.read_excel('Flyzy Flight Cancellation (3).xlsx')
```

[3]:
```
#EXPLORING THE DATA
cancelled_3.head()
```

[3]:

	Flight ID	Airline	Flight_Distance	Origin_Airport	Destination_Airport	Scheduled_Departure_Time	Day_of_Week	Month	Airplane_Type	Weather_Score	Previous_Flight_Delay_Minutes	Airline_Rating	Passenger_Load	Flight_Cancelled
0	7319483	Airline D	475	Airport 3	Airport 2	4	6	1	Type C	0.225122	5.0	2.151974	0.477202	0
1	4791965	Airline E	538	Airport 5	Airport 4	12	1	6	Type B	0.060346	68.0	1.600779	0.159718	1
2	2991718	Airline C	565	Airport 1	Airport 2	17	3	9	Type C	0.093920	18.0	4.406848	0.256803	0
3	4220106	Airline E	658	Airport 5	Airport 3	1	1	8	Type B	0.656750	13.0	0.998757	0.504077	
4	2263008	Airline E	566	Airport 2	Airport 2	19	7	12	Type E	0.505211	4.0	3.806206	0.019638	

```
3                  1
4                  0
```

[4]:
```
cancelled_3.tail()
```

[4]:
```
      Flight ID    Airline  Flight_Distance Origin_Airport  \
2995    1265781  Airline D              395      Airport 2
2996    5440150  Airline E              547      Airport 1
2997     779080  Airline C              461      Airport 1
2998    4044431  Airline B              464      Airport 3
2999    2806578  Airline A              369      Airport 1

     Destination_Airport  Scheduled_Departure_Time  Day_of_Week  Month  \
2995           Airport 3                        0            6      1
2996           Airport 4                       22            4      7
2997           Airport 3                        8            3      1
2998           Airport 3                        5            5      3
2999           Airport 2                        1            1     10

     Airplane_Type  Weather_Score  Previous_Flight_Delay_Minutes  \
2995        Type B       0.190018                        1.00000
2996        Type E       0.719271                       91.00000
2997        Type B       0.458724                        3.00000
2998        Type E       0.443373                       46.00000
2999        Type A       0.704563                       18.66667

      Airline_Rating  Passenger_Load  Flight_Cancelled
2995        2.451216        0.283440                 1
2996        0.027039        0.665294                 1
2997        1.131315        0.991307                 0
2998        0.968651        0.254808                 1
2999        1.879411        0.532486                 1
```

[5]:
```
cancelled_3.columns
```

[5]:
```
Index(['Flight ID', 'Airline', 'Flight_Distance', 'Origin_Airport',
       'Destination_Airport', 'Scheduled_Departure_Time', 'Day_of_Week',
       'Month', 'Airplane_Type', 'Weather_Score',
       'Previous_Flight_Delay_Minutes', 'Airline_Rating', 'Passenger_Load',
       'Flight_Cancelled'],
      dtype='object')
```

[6]:
```
cancelled_3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 14 columns):
```

```
 #   Column                          Non-Null Count  Dtype
---  ------                          --------------  -----
 0   Flight ID                       3000 non-null   int64
 1   Airline                         3000 non-null   object
 2   Flight_Distance                 3000 non-null   int64
 3   Origin_Airport                  3000 non-null   object
 4   Destination_Airport             3000 non-null   object
 5   Scheduled_Departure_Time        3000 non-null   int64
 6   Day_of_Week                     3000 non-null   int64
 7   Month                           3000 non-null   int64
 8   Airplane_Type                   3000 non-null   object
 9   Weather_Score                   3000 non-null   float64
 10  Previous_Flight_Delay_Minutes   3000 non-null   float64
 11  Airline_Rating                  3000 non-null   float64
 12  Passenger_Load                  3000 non-null   float64
 13  Flight_Cancelled                3000 non-null   int64
dtypes: float64(4), int64(6), object(4)
memory usage: 328.2+ KB
```

[8]:
```
#converting all null values to numeric
from sklearn.preprocessing import LabelEncoder
label_encoder=LabelEncoder()
```

[9]:
```
cancelled_3['Airline']=label_encoder.fit_transform(cancelled_3['Airline'])
cancelled_3['Origin_Airport']=label_encoder.
 ↪fit_transform(cancelled_3['Origin_Airport'])
cancelled_3['Destination_Airport']=label_encoder.
 ↪fit_transform(cancelled_3['Destination_Airport'])
cancelled_3['Airplane_Type']=label_encoder.
 ↪fit_transform(cancelled_3['Airplane_Type'])
```

[12]:
```
#all null values are now numeric
cancelled_3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 14 columns):
 #   Column                          Non-Null Count  Dtype
---  ------                          --------------  -----
 0   Flight ID                       3000 non-null   int64
 1   Airline                         3000 non-null   int64
 2   Flight_Distance                 3000 non-null   int64
 3   Origin_Airport                  3000 non-null   int64
 4   Destination_Airport             3000 non-null   int64
 5   Scheduled_Departure_Time        3000 non-null   int64
 6   Day_of_Week                     3000 non-null   int64
 7   Month                           3000 non-null   int64
 8   Airplane_Type                   3000 non-null   int64
```

```
 9   Weather_Score                   3000 non-null   float64
 10  Previous_Flight_Delay_Minutes   3000 non-null   float64
 11  Airline_Rating                  3000 non-null   float64
 12  Passenger_Load                  3000 non-null   float64
 13  Flight_Cancelled                3000 non-null   int64
dtypes: float64(4), int64(10)
memory usage: 328.2 KB
```

[13]:
```
cancelled_3.isnull().sum()
```

[13]:
```
Flight ID                       0
Airline                         0
Flight_Distance                 0
Origin_Airport                  0
Destination_Airport             0
Scheduled_Departure_Time        0
Day_of_Week                     0
Month                           0
Airplane_Type                   0
Weather_Score                   0
Previous_Flight_Delay_Minutes   0
Airline_Rating                  0
Passenger_Load                  0
Flight_Cancelled                0
dtype: int64
```

[14]:
```
cancelled_3.describe()
```

[14]:

	Flight ID	Airline	Flight_Distance	Origin_Airport	\
count	3.000000e+03	3000.000000	3000.000000	3000.000000	
mean	4.997429e+06	1.567333	498.909333	1.631667	
std	2.868139e+06	1.513350	98.892266	1.499805	
min	3.681000e+03	0.000000	138.000000	0.000000	
25%	2.520313e+06	0.000000	431.000000	0.000000	
50%	5.073096e+06	1.000000	497.000000	1.000000	
75%	7.462026e+06	3.000000	566.000000	3.000000	
max	9.999011e+06	4.000000	864.000000	4.000000	

	Destination_Airport	Scheduled_Departure_Time	Day_of_Week	\
count	3000.000000	3000.000000	3000.000000	
mean	0.911667	11.435000	3.963000	
std	1.147012	6.899298	2.016346	
min	0.000000	0.000000	1.000000	
25%	0.000000	6.000000	2.000000	
50%	0.000000	12.000000	4.000000	
75%	2.000000	17.000000	6.000000	
max	3.000000	23.000000	7.000000	

```
              Month  Airplane_Type  Weather_Score  \
count  3000.000000    3000.000000    3000.000000
mean      6.381000       1.582000       0.524023
std       3.473979       1.515049       0.290694
min       1.000000       0.000000       0.000965
25%       3.000000       0.000000       0.278011
50%       6.000000       1.000000       0.522180
75%       9.000000       3.000000       0.776323
max      12.000000       4.000000       1.099246

       Previous_Flight_Delay_Minutes  Airline_Rating  Passenger_Load  \
count                    3000.000000     3000.000000     3000.000000
mean                       26.793383        2.317439        0.515885
std                        27.874733        1.430386        0.295634
min                         0.000000        0.000103        0.001039
25%                         7.000000        1.092902        0.265793
50%                        18.000000        2.126614        0.517175
75%                        38.000000        3.525746        0.770370
max                       259.000000        5.189038        1.123559

       Flight_Cancelled
count       3000.000000
mean           0.690667
std            0.462296
min            0.000000
25%            0.000000
50%            1.000000
75%            1.000000
max            1.000000
```

[15]:
```
cancelled_3.shape
```

[15]:
```
(3000, 14)
```

[16]:
```
sns.pairplot(cancelled_3,hue='Flight_Cancelled')
```

[16]:
```
<seaborn.axisgrid.PairGrid at 0x7f00e9db6c50>
```

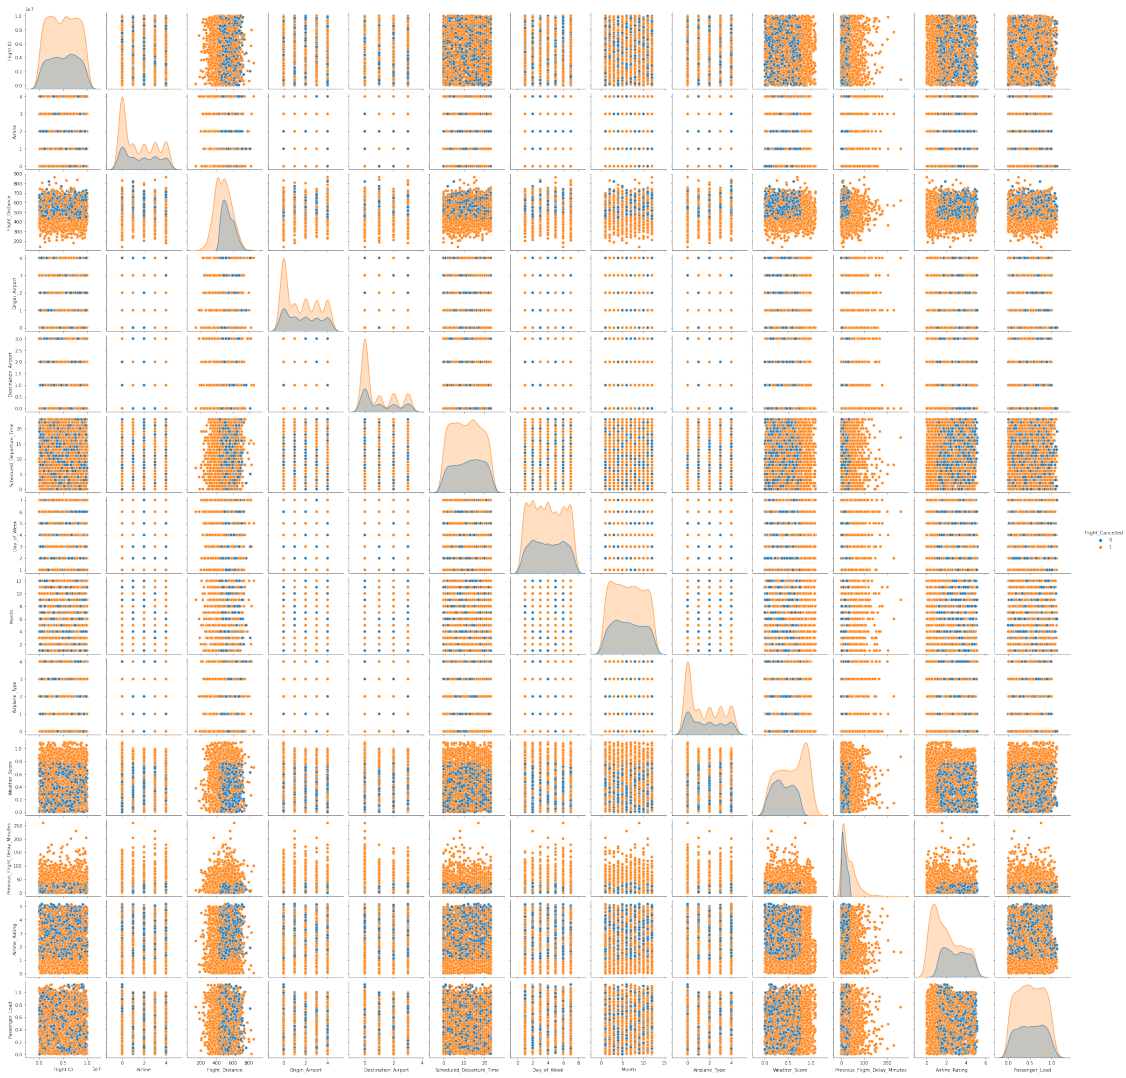

[28]:
```
#Relationship analysis
corelation=cancelled_3.corr()
```

[31]:
```
sns.heatmap(corelation, xticklabels=corelation.columns, yticklabels=corelation.
columns
,annot=True)
```

[31]: <AxesSubplot: >

[55]:
```
x=cancelled_3[['Flight ID', 'Airline','Flight_Distance', 'Origin_Airport',␣
 ↪'Destination_Airport','Scheduled_Departure_Time','Day_of_Week',
'Month', 'Airplane_Type', 'Weather_Score',␣
 ↪'Previous_Flight_Delay_Minutes','Airline_Rating','Passenger_Load']]
y=cancelled_3['Flight_Cancelled']
```

[34]:
```
correlation=np.corrcoef(x['Flight ID'], y)[0, 1]
print("Correlation coefficient:", correlation)
```

```
Correlation coefficient: -0.0091005441027218
```

[35]:
```
correlation=np.corrcoef(x['Airline'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: -0.0579152207425434
```

[36]:
```
correlation=np.corrcoef(x['Origin_Airport'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: -0.049925451318859365
```

[37]:
```
correlation=np.corrcoef(x['Destination_Airport'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: -0.06475308532828986
```

[38]:
```
correlation=np.corrcoef(x['Destination_Airport'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: -0.06475308532828986
```

[39]:
```
correlation=np.corrcoef(x['Scheduled_Departure_Time'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: -0.043732799217209566
```

[40]:
```
correlation=np.corrcoef(x['Day_of_Week'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: -0.008705376908751991
```

[41]:
```
correlation=np.corrcoef(x['Month'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: -0.004242162010093053
```

[42]:
```
correlation=np.corrcoef(x['Airplane_Type'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: -0.04994233077062403
```

[43]:
```
correlation=np.corrcoef(x['Weather_Score'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: 0.30576162505311555
```

[44]:
```
correlation=np.corrcoef(x['Previous_Flight_Delay_Minutes'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: 0.302804640547877
```

[45]:
```
correlation=np.corrcoef(x['Airline_Rating'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: -0.31409863751544576
```

[46]:
```
correlation=np.corrcoef(x['Passenger_Load'],y)[0,1]
print("Correlation coefficient:",correlation)
```

```
Correlation coefficient: -0.008319756091970014
```

[47]:
```
from sklearn.linear_model import LinearRegression
```

[48]:
```
model=LinearRegression()
model.fit(x, y)
print("Intercept:", model.intercept_)
print("Coefficients:", model.coef_)
```

```
Intercept: 1.1322946462844234
Coefficients: [-7.78946167e-11 -1.00440931e-03 -1.28197407e-03  3.22053796e-03
 -1.33395687e-03 -2.33343775e-04 -2.62244908e-03  2.19967335e-03
  3.11880877e-03  4.85104135e-01  5.16096743e-03 -8.89749801e-02
  7.14363056e-03]
```

[49]:
```
#EVALUATION  R-SQUARED.
from sklearn.metrics import r2_score
```

[50]:
```
y_pred=model.predict(x)
r2=r2_score(y, y_pred)
print("R-squared:", r2)
```

```
R-squared: 0.3495804329516544
```

[51]:
```
dataset_2=pd.read_excel("Flyzy Flight Cancellation.xlsx")
```

[62]:
```
x=cancelled_3.drop(columns=['Flight ID', 'Airline','Flight_Distance',␣
 ↪'Origin_Airport','Origin_Airport',␣
 ↪'Destination_Airport','Scheduled_Departure_Time','Day_of_Week',
'Month', 'Airplane_Type', 'Weather_Score',␣
 ↪'Previous_Flight_Delay_Minutes','Airline_Rating','Passenger_Load'])
y=cancelled_3['Flight_Cancelled']
```

[73]:
```
#model feeding.
model=LinearRegression()
model.fit(x_train, y_train)
```

[73]:
```
LinearRegression()
```

[74]:
```
y_pred=model.predict(x_test)
```

[75]:
```
r_squared=r2_score(y_test, y_pred)
print("R-squared:", r_squared)
```

```
R-squared: 1.0
```